CEC460 Final Project: PCI Express

David

1. Identify the entities that communicate with one another using this protocol. In most cases there are two entities with one acting as the client and the other as the server, or caller and callee. But in some protocols there may be three kinds of entities with a third entity acting as a manager, facilitator, or go-between.

   Peripheral component interconnect express (PCIe) is a full duplex serial communication interface often used in communication between a CPU and some high speed peripheral. Devices often connected using PCIe include graphics cards, high-speed Ethernet/Wi-Fi, and modern hard drives. The first revision of PCIe was designed in 2003 to replace the existing parallel buses PCI (and its variations) and AGP that were widely used but starting to run up against bottlenecks preventing further increases in bandwidth. As a parallel bus increases its operating frequency, jitter and skew become significant problems to overcome. The old conventional PCI, which was used as a general connection between a computer motherboard and various peripherals, used a 32 or 64-bit parallel bus connection between devices. Additionally, the PCI bus was a true bus, with almost all wire connections shared by all devices using the bus. This meant that for a bus connection involving many devices, the electrical connection might need to travel a long distance and have high capacitance, which make increasing the bus's speed past the MHz range difficult. The parallel bus also needed many traces on the motherboard to route all required signals. AGP was another parallel interface used for graphics cards, subject to most of the same concerns as conventional PCI.

   For these reasons, the PCI special interest group (PCI-SIG) left parallel buses behind in favor of serial interfaces, which replace the problems discussed above with other difficulties that are challenging but possible to overcome. PCIe started to replace all uses of PCI and AGP in 2004, and it is used almost exclusively as the connection to CPU peripherals on modern motherboards.

Layered Architecture

PCIe implements an architecture of abstraction layers similar to the OSI network model. While the 5-layer OSI model has 4 supporting layers below the application layer, PCIe has a simple enough connection topology that it only uses three: transaction, data link, and physical layers. The 3 PCIe layers perform functions comparable to the lowest 4 layers of the OSI model, without the need for a network layer that handles routing between links in a complicated network topology. The PCIe layers are shown in the figure below.

The uppermost layer of the PCIe stack is the transaction layer which fulfils a role like OSI's transport layer. The transaction layer handles information being sent through the PCIe link in a unit of information called a transaction layer packet (TLP), which carries different information depending on the type of transaction it is being used for. The four types of transactions supported by PCIe are memory, I/O, configuration, and message transactions, of which the three former were inherited from similar packet types in the PCI protocol. Message transactions, which were not part of earlier parallel interface standards, carry metadata about the link, such as interrupt requests, that were formerly handled by discrete wire connections in PCI.

The next PCIe layer is the data link layer, analogous to the OSI layer of the same name. The PCIe data link layer is responsible for managing links between PCIe devices and ensuring reliable data delivery. The information that the data link layer adds to TLPs it receives from the transaction layer include a sequence number for uniquely identifying packets at reception and a CRC for ensuring that data being read from the link is valid. The data link layer often needs to communicate with the data link layer of another device for link management even when no transaction layer information needs to be sent. To meet this end, the layer produces a different kind of packet, called a data link layer packet (DLLP), which is distinct from a TLP.

The lowest layer on the PCIe stack is the physical layer, which pertains to bits of a message from the layers above being sent along a hardware connection to be received by the physical layer on the other end of a link.
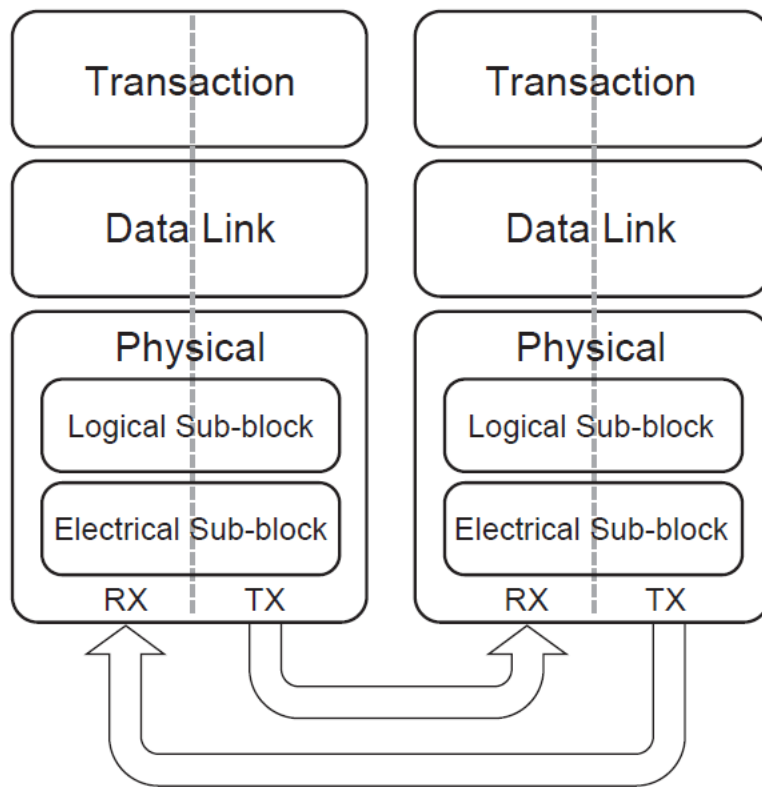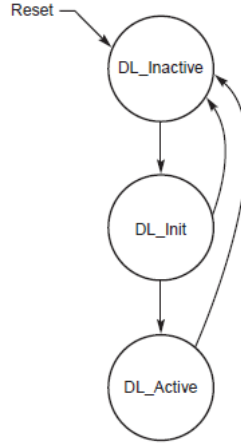
Figure 1: PCIe Layers

Figure 2: Data link control and management state machine diagram

The PCIe interface is capable of transferring information at incredibly high rates, in the GB/s range, so the physical connection needs to operate at a very high speed. Standard wire connections are subject to increased noise and power consumption when very high clock rates are involved, so to combat these negative effects PCIe uses differential pairs of electrical signals to transmit information. At this layer, the bits are also encoded to ensure that clock and synchronization information can be properly extracted from the signal. Though PCIe uses a shared reference clock to ensure that devices connected by a link are operating on the same frequency, the data transitions are not synchronized to any edge transition of this clock. Instead, each device has a PLL monitoring bit transitions on the data lines that reconstructs the serial synchronization based on when bit transitions occur. To ensure that enough bit transitions occur to extract necessary synchronization, the PHY encodes each 8 bit word using 10 bits in an 8b/10b format that enforces DC-balance and enough signal transitions to allow clock recovery. This means that every PCIe transaction has only 80% efficiency of bits sent, that is before version 3.0 which replaced the 8b/10b encoding with a 128b/130b encoding with scrambling.

2. Describe how devices identify themselves on the network, and become known to other devices on the network. If initial connection involves a handshake, describe the handshake. If identities are given, describe how the identities are assured to be unique.

   The PCIe link is managed by the data link layer using a finite state machine called the data link control and management state machine (DLCMSM). The DLCMSM can be in one of three states: DL_Inactive, DL_Init, or DL_Active, wherein the link has as status output either DL_Up or DL_Down. The names of the states follow what intuition implies, but they will be discussed in further detail below. The transitions between states in the DLCMSM are shown in the diagram below.

   In the DL_Inactive state, which is entered upon the device powering up, the link is non-operational or nothing is connected. The link can move to the DL_Init state if the data link layer receives a signal from the physical layer that the physical link is active. In this state the status is DL_Down.

   While in the DL_Init state, the link works at initializing the flow control (FC) to default values. This process has its own two states responsible for configuring enabled virtual channels for use. While in the first FC initialization state, the status is DL_Down, but upon moving to the second the status is finally DL_Up. If at any point during the FC init process the physical layer reports the that the physical link is down, the DLCMSM moves back to the DL_Inactive state. If the FC init completes successfully without the physical link going down, the DLCMSM moves on to the DL_Active state. Part of FC initialization involves configuring multiple devices on a switch with different traffic classes and virtual channels to differentiate between communication involving certain components but not others. This flow control information is persistent through the entire link for every message sent from one device to another. Initialization is primarily performed by the physical layer, which negotiates the number of

lanes and data rate shared by communicating devices.

Once every initialization step is complete, the link enters the DL_Active state, the normal operating mode for the PCIe link. In this state, TLPs and DLLPs can be transmitted over the physical layer as necessary and the transaction layer link only goes down if the physical link goes down, through device removal or failure.

3. Discuss the issue of traffic and congestion. Describe how the network and protocol addresses the problem of either too many devices trying to use the network at the same time or trying to push too much data through the network. Do devices take turns (how do they know?). Do devices tell get told to stop or wait? Do devices get told to slow down or speed up? How are control signals separated from data?

PCIe is a point-to-point connection, so a single "bus" can only natively connect two devices. The PCIe topology includes at its base a root complex, the interface connecting PCIe links to the CPU through a parallel interface, typically on-chip. The root complex generally allows the CPU to have access to multiple PCIe links, but each link can only connect to one component at the other end. More devices can be connected using a switch, which allows a single link from the root complex to connect to multiple other devices, emulating a true bus. PCIe supports traffic control using what it calls traffic classes and virtual channels. The upstream (towards CPU) end of a switch can only have one physical PCIe port, so if the downstream end has multiple links to different devices they must multiplex the connection to the root complex. Each link to the downstream end of a switch is assigned a traffic class which can be thought of as a priority of the traffic associated with a given device. Discrete links connected to a PCIe channel are identified in part by their traffic class, contributing to something similar to an Ethernet MAC address, but on a much smaller and more dynamic scale.

A major part of PCIe is that a physical connector can be used with any PCIe device of any capability, including vastly different bandwidth and number of lanes. For a point-to-point full duplex architecture, there is no danger of a collision between two devices attempting to transmit at the same time, as only one device is capable of writing to each data line. When a switch is used to make the architecture more bus-like, the switch is responsible for regulating which virtual channels have access to which lanes at a time, and multiplexes between different low-bandwidth devices trying to talk on the same bus.

Sean

4. Discuss the issues of security.

a. Does the protocol address the issue of privacy? If so, how? If not, what risks might exist?

b. Does the protocol protect against malicious interlopers on the network? If so how? If not, how might the network be exploited or compromised by someone who could connect a malicious device to the network?

The PCIe protocol does not have features that directly address the issues of security. However, given the physical characteristics of the protocol, limited wire length for data transmissions and high operating frequencies. The hardware is difficult to tap from a distance. Should one want to access data being transmitted directly from PCIe communications, they would need direct access to the PCIe slots that facilitate communication between two devices.

c. How does the protocol recover after a failure, i.e. after a power outage or some other breakdown?

5. Is efficiency a concern? Does the protocol have features to optimize data transmission rates? What does it do to control or adjust speeds? What if different devices have different speed capabilities?

Efficiency is a large concern for the PCIe protocol. One of the main reasons why you would want to use the PCIe protocol is for fast communication rates. The latest version of PCIe (version 5.0 at the time of writing this documentation) can obtain a maximum throughput of 64 GB/s using 16 lanes. One contributor that allows PCIe to achieve these fast transmission rates is each device has its own serial medium connecting to the root complex or switch. The purpose of the switch is to allow multiplex

Figure 3: Legacy Medium Configuration

connections between the PCIe peripherals on a single link with the root complex or a switch. This varies drastically when compared to the PCI protocol which has all devices sharing one medium. The PCI medium connections significantly limits the rate at which data can be transferred since only one device can transmit data at a time. Images depicting these differences in medium connection between the PCI and PCIe protocol are as follows.

It is worthy to note the PCIe protocol can mock the bus configuration of PCI using switches. This allows for software backwards compatibility between the two protocols.

In order to ensure that bits are sent and received at the same rate between two devices an addition 2 bits are inserted for each 8. These added bits allow the phase locked loop (an additional piece of hardware) to measure the frequency at which signals are transmitted and synchronize devices.

6. How does the protocol address communication failures? What kinds of failures are addressed? Suggest two kinds of failures that are not addressed, and discuss what could happen to senders and receivers in a failure of that kind.

The PCIe protocol accounts for transmission failures by ensuring that the stream of bits that were sent from a device, were the same as the ones received. This is achieved via Cyclic Redundancy Checks (CRC). In the case that the CRC identifies an error, a request will be sent to the sending device to resend those packets.

One unaccounted possibility for error is in the transaction layer (the transaction layer is one level of abstraction above the data link layer and is specific to the PCIe protocol). The transaction layer has its own header and is as follows.

The purpose of the Transaction Layer Packet (TLP) header is to interpret read and write instructions. Once the device has fulfilled the request of the read or write instruction, a new packet is sent back to the device that initiated the request. The PCIe protocol is unclear as to what steps are taken when more than one device share the same Requester ID.

Figure 4: Express Medium Configuration

Figure 5: PCIe TLP header